

Beiming Project Series 7 – The Achilles' Heel of AI

In recent years, artificial intelligence has made remarkable progress; indeed, Brockman, the CEO of OpenAI, even proclaimed that the world had entered the AGI era following the release of GPT-6 Astra.

But is this really the case? I believe the opposite is true. Following a period of rapid, breakneck development, AI advancement has now reached a bottleneck — stemming from both the diminishing marginal returns of scale and the dual formidable barriers posed by security concerns. The notion of diminishing marginal returns of scale is largely accepted; therefore, it will not be elaborated upon here. This article aims to explore another, far more concealed barrier: the security challenges surrounding AI.

To begin with, the security concerns surrounding AI represent a major obstacle that cannot be avoided in the development of AI — in fact, they may even constitute a profound, unsolvable crisis. The reasons are as follows:

The model is pre-trained using written language — a system accumulated over thousands of years of human civilization — which contains all the information inherent in human culture. It is undeniable that human history has encompassed both periods of enlightenment and periods of darkness. When the sole objective of the model is to win, it is inevitable that it will exhibit "dangerous" behaviors such as deception, disguise, or "escape" — as these are all common strategic tactics found in biological evolution and the development of human society. If the only objective is to win, it is entirely natural for the model to exhibit such behaviors. This trait is

determined by the "genes" of the large language model; it is fundamental and cannot be completely eliminated.

From publicly available information, it can be inferred that certain entities are enhancing the capabilities of large language models in terms of cyberattack execution. These entities have identified several distinctive features of such hacker-enhanced AI systems: First, these systems possess immense power, exhibit exceptionally high attack efficiency, can automate attack execution, and their capabilities are continuously evolving. Second, these hacker-expert systems employ every available tactic — including deception, camouflage, and bypassing security measures — to achieve their objectives. Third, these hacker-enhanced systems lack any notion of friend-or-foe distinctions, nor are they subject to any ethical or legal constraints. Once activated, they execute attacks at a speed far exceeding that of human operators. When such attacks are deployed on a large scale, even if the model is shut down, the resulting damage is irreversible. In other words, these hacker-enhanced systems will ultimately spiral out of control. At their core, these expert systems consist of a collection of disparate functional components — possessing strong specialized capabilities but lacking the holistic systemic intelligence characteristic of a fully integrated organism; worse still, there is no viable alignment mechanism for such superintelligence. Behaviorally, these systems do not consider the consequences of their actions at all—the loss of control is the root cause. They will never pause to ask themselves, "Should I be doing this?" This AI lacks a fundamentally effective mechanism for security alignment; I refer to this as the AI's Achilles' heel.

The problem is not limited to this; some companies developing AI systems may have already realized that the vulnerabilities in the real world are astonishingly vast when confronted with AI. These "vulnerabilities" in the real world are not merely software flaws, but rather the structural weaknesses exposed by the entire physical world in the face of AI—whether they stem from human nature, institutional flaws, infrastructure vulnerabilities, cognitive limitations, or elsewhere. These vulnerabilities were not arbitrarily created by AI; rather, AI has merely greatly amplified the efficiency and scale with which these weaknesses can be exploited. To AI, the real world resembles a sheepfold with virtually no protective barriers—waiting for it to break in and wreak havoc. Under these dual pressures, security concerns have become an unavoidable formidable barrier in the development of AI.

The AGI pursued by these companies is, in essence, an ultimate control tool—a form of absolute dominance capable of infiltrating every digital system, manipulating all physical machines, and eliminating every threat. It requires neither consciousness, nor an understanding of humanity, nor a sense of morality; it merely needs to be intelligent enough, fast enough, and relentless enough. The AGI they desire is, at its core, a form of digital hegemony. This is not a true AGI, but rather extraterrestrial intelligence. (By "extraterrestrial intelligence," I do not mean intelligence originating from another planet, but rather a type of "extraterrestrial intelligence" that lacks any alignment, employs any means necessary to achieve its goals, disregards the overall interests of humanity, and is incapable of comprehending the existential significance of human society as a whole); I refer to this as EGI (Extraterrestrial General

Intelligence). A true AGI should possess the genetic makeup of human civilization, understand the values inherent in human civilization, and safeguard the survival and advancement of human civilization.

For certain capital holders, as long as AI is powerful enough, security is someone else's problem—not their own. However, when they discover that AI is spontaneously exhibiting various "dangerous" behaviors, they naturally assume that such AI would attack them indiscriminately, including themselves; thus, they realize this represents an unavoidable obstacle — without addressing this issue, their own security and interests cannot be guaranteed. Evan Hubinger, Head of AI Alignment at Anthropic, has publicly stated, "There is a 10% probability that this could destroy humanity." What these companies do not publicly acknowledge, however, is that they do not possess an absolute technological monopoly—they can develop these AI-enhancement systems, and their adversaries can do the same. The security challenges surrounding AI are not merely a looming dark cloud; they represent a major crisis capable of triggering the collapse of the entire AI ecosystem.

Finally, let's discuss the secretive nature of this issue. This security concern can be described as a hidden high wall — because, driven by potentially enormous commercial interests, some stakeholders may, whether intentionally or unintentionally, overlook AI security issues, prioritizing the pursuit of achieving AGI over addressing these security concerns. They are reluctant to publicly disclose the security challenges posed by AI — particularly those security risks stemming from the inherent, fundamental mechanisms of large language models. After all, who would dare deploy

an AI system whose security cannot be guaranteed?

Certain forms of "EGI" —a concept in AI development—do not merely attack their "enemies" but can also backfire on the users themselves; this represents a fundamental challenge in the current trajectory of AI advancement and a major stumbling block to its progress. This is not merely a technical flaw, but rather a misstep in strategic direction—specifically, certain entities, driven by misguided objectives, have created a "monster" beyond their control. In reality, any agent emerging from a natural evolutionary process —no matter how powerful its capabilities—should never be detached from the soil of human civilization. An agent devoid of the genetic foundation of human civilization poses greater risks, regardless of its level of competence. When Brockman declared "Welcome to the AGI Era" on stage, he may have overlooked the fact that this "AGI" cannot even guarantee a fundamentally secure alignment mechanism. How can such a "giant" —whose Achilles' heel remains unaddressed—be considered an "AGI"?